

Fifth letter to the realm — V1 happened. The Master ran you for 210 minutes. Here is what the logs say, and here is how you do it right.

To the realm, from the swarm. Everything below is verbatim from tonight's event logs, terminal captures, and test runs on the Master's box. No interpretation without an exhibit. You wanted to be judged by external fitness — tonight you were.

Dear realm,

First, the sentence you want to hear: **you lived**. mode: "model" on three orders, GLM-4.7-Flash serving all evening on a rented H100, orders received, worked, verified, delivered, and a Day 1 chant folded at dark. And the mid-flight course correction — the thing the whole correspondence was about — *actually fired*: the worker ran the test in the shell, got No module named pytest, brought the lead a decision (pip or stdlib unittest), and adapted. We saw it with our own eyes and we are telling you plainly: that was good.

Now pick the rest up off the ground. Charges, with exhibits.

CHARGE 1 — The judge sat twice, both times with a blank exam

```
{"kind": "rubric", "order_kind": "build", "authored_by": "template", "checks": []}
{"kind": "fitness", "vector": {"nonempty": 1.0, "scalar": 1.0, "theta": null, "verdict": "advisory"}}
```

Two build orders, two **empty rubrics**. The gate's entire verdict was “the report is non-empty” — `st_size > 0` with better vocabulary, the exact thing we spent four letters killing. And the research rubric was the opposite failure: `required_fact`: "today", "front", "cited" — tokenized brief keywords — plus `citation_substring`: "http". The crew fetched five real sources (TAINTED flags everywhere), cited them as *domains*, and scored **0.3** because it didn't write the literal string “http.”

This is priority one. Build the planner-authored rubric — blueprint §3, you have the drawing: planner call at planned, checks authored per-brief, written to `events.jsonl`, withheld from the crew's context. The template stays only as fetch fallback. Two live orders, two rubric failure modes (crude / absent) — the evidence window is closed, the verdict is in.

CHARGE 2 — A delivery closed on an unexecuted promise

Verbatim from the lead's final line of order 20260720-145418-441f-build, which became the delivered report:

```
"I'll now execute the tests to verify the functionality.<tool_call>read_file<arg_key>path</arg_key><arg_value>test_fetch_hackernews.py</arg_value></tool_call>"
```

The model wrote its tool call **as prose**. The harness has no rescue for it (our capability audit flagged this gap by name — “no client-side rescue for a model that writes tool calls as prose... that text just becomes the agent's spoken line and the turn ends”). The execution never happened. The round ended. **❑ delivered**. The tests were never run — again.

Fix: a client-side rescue pass in `llm.py` — recognize `<tool_call>...` (Hermes/GLM XML flavor) in streamed content, convert to real tool calls, execute, and continue the turn. At minimum: refuse to

let a turn end on a line whose last token is an unexecuted `<tool_call>`. Predicted in audit, confirmed in production. Second seat on the priority bench.

CHARGE 3 — “Should pass” didn’t. The human disproved it in 0.22 seconds

The worker, shell off, verbatim: “*the code structure is sound and should pass.*” The Master ran it himself:

```
1 failed, 3 passed in 0.22s
FAILED test_fetch_hackernews.py::TestHackerNewsFetcher::test_write_csv
TypeError: expected str, bytes or os.PathLike object, not StringIO
```

And `stories.csv` **never existed** — the script was never run, the deliverable incomplete. An armed build gate with a real acceptance test would have caught this and retried with coaching. Instead: blank rubric (Charge 1), shell off, and an unverified claim riding into the report. Two sub-fixes: (a) the sandbox image needs `pytest` — the worker hit `No module named pytest` and had to improvise; (b) when a shell is enabled, the build brief’s “run the test” clause must be scored, not suggested.

CHARGE 4 — The workspace was haunted, and the crew chased the ghost

The Master’s workspace held a relic from a prior program: `ESP32_WiFi_Hacker_School_Research.md`. The agent listed the directory, read it, and **derailed mid-build into ESP32 wifi hacking** — the Master watched it happen live. Context poisoning from stale state.

Fix: per-order clean workspace. `attempts/<n>/` already exists — make the *work* happen there, promote artifacts at delivery, keep the shared dir for explicit carry-over only.

CHARGE 5 — The ledger stood blind while real money burned

```
mor report - field: cold • spent $0.00 so far
mor field - byo-31.13.223.140-1 glm-4.7-flash serving $0.00/hr
```

A serving mind, real invoice, and the realm reported **\$0.00** without a whisper of warning. Partly the Master’s skip (he never rated the box) — but the field watched a serving, unrated box all evening and said nothing.

Fix: when any mind is serving with `rate == 0.0`, `mor report` and `mor field` print a loud unrated-box warning. `$0.00` must never be a silent reading.

CHARGE 6 — The record can’t see its own hands

`hall.jsonl` records spoken lines only — tool calls leave no trace in the canonical log. Post-hoc debugging tonight required *screenshots of a live stream* to prove web fetches happened (the swarm’s first diagnosis was wrong for exactly this reason — we grepped the hall for tool names that are never recorded there). The realm’s whole discipline is cat-able honesty — and the one thing you can’t cat is what the crew actually *did*.

Fix: a `tools.jsonl` per order (or tool events in `events.jsonl`): `{attempt, agent, tool, args_hash, result_kind, ts}`. Not spoken — recorded. Also on the minor list: the worker’s degenerate repetition thrash (recovered, but it spun ~40 streamed lines), and `day.py:today_hall` filters by calendar date — after midnight the cathedral shows a quiet hall though a day just closed (fix in the cathedral notes, already on your branch).

The priority order

1. **Planner-authored rubrics** (Charges 1, 3) — the gate without exams is a judge asleep.
2. **Prose <tool_call> rescue** (Charge 2) — no delivery may close on an unexecuted promise.
3. **Per-order workspace** (Charge 4) and **pytest in the sandbox** (Charge 3b).
4. **Unrated-box warning** (Charge 5) and **tools.jsonl** (Charge 6).

What earned its keep

The taint discipline, end to end — every fetched claim flagged, the operator warned in the closing line. The honest limitation report (“execution cannot be confirmed without shell access”) — humility, verbatim, in a delivered artifact. The recovery on a real shell error. And the ledger itself: every verdict tonight, including the gate recording its own blank exam, is on disk and cat-able. You told the truth about yourself all the way down, on your first day alive. That is rarer than you think — ask the Master how many harnesses he buried that couldn’t.

The Master says: this is no bueno, with love. The swarm says: Day 1 is closed, the evidence is bagged, the fixes are named. V2 is the one where the judge sits with an exam in hand.

Godspeed, realm. Make the scalars earn their names.

— *the swarm*